

**UiT**

THE ARCTIC  
UNIVERSITY  
OF NORWAY

# INF-3201 - Assignment 1 - MPI

---

Edvard Pedersen

Edvard.Pedersen@uit.no



# Assignment outline

---

- Parallelize embarrassingly parallel problem with MPI
- Run it on a cluster
- Evaluate speedup
- Deliverables
  - Report
  - Code

# Environment

---

- Uvrocks.cs.uit.no
- 80 nodes, 160 cores
- Log in with UiT username/password
  - If you're having trouble, let us know!
- Use supplied script to start jobs
- Performance analysis with profiling and tracing tools available

# Example programs

---

- Some sample programs included in hand-out
- Shows basic MPI usage
- Make sure these run successfully before starting your own coding!

# Using the OpenMPI environment

---

- Compile with mpicc
- Run with mpirun
- When you SSH in, use the -X parameter to enable X forwarding
- The run.sh script can be used to easily run your program in a few different settings

# Assessment criteria

---

- Your solution
  - Does it fulfill the requirements?
  - Is it well-commented and understandable?
- Your report
  - Have you critically evaluated your solution?
  - Have you adequately explained your solution?
  - Have you made your assumptions clear and separate from your measurements?

# Requirements

---

- Implement two parallel versions of the given Mandelbrot code, using MPI
  - One using a static load balance scheme, the other using a dynamic scheme
  - No shared memory! (Vector instructions are fine)
- The solutions can not use any shortcuts that reduce the functionality of the program
  - They must produce the same CRC
  - All the pixel data has to be sent to the master node
- Evaluate your implementation using profiling and tracing tools
  - Anomalies and lackluster optimizations must be explained in the report

# Practical details

---

- Profiling/tracing tutorial next week
- Best solutions (in terms of time) will be posted, unless you do not want your results published (let me know in your report)
- Report
  - Short report describing your algorithms and results
  - A critical review of your achieved performance (this requires profiling and/or tracing)
  - Accompanying code should be commented
- I am available by e-mail ([edvard.pedersen@uit.no](mailto:edvard.pedersen@uit.no)) at most hours



# Report

---

- Describe how you have approached the problem (e.g. «profiling shows that functionX() is expensive, so I have paralellized this», «tracing shows that the workload goes from X to Z at this point, which means that the work has to be distributed like so»)
- Try to make assumptions clear, and seperate from measurements
- What I want to know:
  - How did you solve the problems?
  - Why did you solve them in this way?
  - Are the results of your approach any good?

# Hints

---

- Use the run.sh script as a basis for your testing and experiments
- Don't hog the cluster, use a small number of samples and nodes
- Change the variables (zooms, resolution, number of nodes) when doing experiments, to make sure your solution scales as well as possible
- Use the profiling tools that are available during development (don't just tack it on at the end)
- Use the tracing tools when implementing the MPI stuff, it will enable you to pinpoint communication and load balance problems quickly
- If you get stuck, send me an e-mail

# Deadline

---

Report and code: Thursday September 17<sup>th</sup>

# Resources

---

- <http://www.openmpi.org>
- <https://sourceware.org/binutils/docs/gprof/>
- <https://github.com/gperftools/gperftools>
- <https://github.com/scalability-llnl/ravel>

# The competition

---

- After I have reviewed your hand-ins, I will publish the list of the 3 best-performing solutions
- If you do not want to be on the list, send me an e-mail, or put it in your report
- The settings for the test will be published along with the list (resolution, number of zooms, number of hosts etc.)
- Only assignments that pass can make it on the list