

UiT

THE ARCTIC
UNIVERSITY
OF NORWAY

INF-3201 - Assignment 3 – GPGPU

Edvard Pedersen
Edvard.Pedersen@uit.no



Assignment outline

- Parallelize the given piece of code with CUDA/OpenCL and multiprocessing (combined)

Environment

- Your own computer (Preferred)
- The lab computers (ifilab110-123) - Possibly problematic
- UVRocks frontend
- Precode in Python
 - But feel free to solve it using whichever language you prefer
 - Precode also shows how to interface to C if you don't want to reimplement all of the code

Example programs

- Included in the handout (src/hello_mandelbrot/):
- Example of using CUDA C (cuda_kernel.cu)
- Example of using PyCUDA to use the CUDA C kernel
- Example of a Python extension written in C
- Example of how to use a C library in Python
- Requirements:
 - Numpy
 - Matplotlib (if you want to look at the images)
 - Pycuda
 - Gcc
 - CUDA Toolkit

Precode

- The precode consists of code to encode and decode messages using XTEA and CBC, as well as a simple mangling algorithm
 - Frontends (encrypt.py and decrypt.py)
 - Main code (precode.py)
- Also has code to attempt to guess the password given a known part of the plaintext

Encoding

- Each block of decoded data consists of two data points
 - Each data point consists of 24 bits of location information and 8 bits of data
- $\text{final_result}[\text{datapoint} \ \& \ \text{location_mask}] = \text{datapoint} \ \& \ \text{data_mask}$
- The final result is random bits, with the secret embedded in a random location
- Each block of the final result (64 bits) encoded with XTEA and CBC
 - Very simple encryption algorithm

Decoding

- Decoding is the same, but in reverse order
- When we don't have a password, we guess by iterating through the printable characters for several levels
 - The number of possible passwords grows incredibly fast
- To test if a password is correct, the data is decoded with the guessed password, then we try to unmangle it
- If we find something we know is in the plaintext in the decoded data, we assume the password is correct

Additional details about the precode

- Four encrypted files included (easy.npy, medium.npy, hard.npy, veryhard.npy)
- All have a secret text which includes the word «secret»
- A very simple GPU implementation decrypts the hard.npy file in ~7 seconds
 - The precode decrypts the easy.npy file in ~26 seconds

If you want to implement it in C (no python)

- Use nvcc to compile (included in the CUDA toolkit)
- <http://docs.nvidia.com/cuda/>
- Alternatively: OpenCL
- <https://www.khronos.org/opencv/>
- <http://developer.amd.com/tools-and-sdks/opencv-zone/amd-accelerated-parallel-processing-app-sdk/>
- <https://software.intel.com/en-us/opencv-code-builder>

Assessment criteria

- Your solution
 - Does it fulfill the requirements?
 - Is it well-commented and understandable?
- Your report
 - Have you critically evaluated your solution?
 - Have you adequately explained your solution?
 - Have you made your assumptions clear and separate from your measurements?

Requirements

- Parallelize the pre-code using GPU and CPU
 - Use multiple cores on the CPU, and GPU, at the same time
 - The CPU parallel implementation can use shared memory or message passing models
- Analyze your solution using the Nvidia Visual Profiler or equivalent for other GPUs
 - Also analyze the CPU parallelization (but this doesn't have to be detailed)

Practical details

- Report
 - Short report describing your algorithms and results, as well as reasons for choosing this approach
 - A critical review of your achieved performance (this requires profiling and/or tracing)
 - Accompanying code should be commented
- I am available by e-mail (edvard.pedersen@uit.no) at most hours

Report

- Describe how you have approached the problem (e.g. «profiling shows that functionX() is expensive, so I have paralellized this», «tracing shows that the workload goes from X to Z at this point, which means that the work has to be distributed like so»)
- Try to make assumptions clear, and seperate from measurements
- What I want to know:
 - How did you solve the problems?
 - Why did you solve them in this way?
 - Are the results of your approach any good?

Deadline

Report and code: Thursday October 29th

Competition

- Another informal competition!
- Let me know if you want to be excluded from the results
- Same procedure as last time
 - Except: I will run the code on a machine with a Intel Xeon E5-1620 and an Nvidia Quadro 600

Resources

- <http://docs.nvidia.com/cuda>
- <https://www.khronos.org/opencl/resources/opencl-tutorials-white-papers-and-how-to-guides>
- <http://mathematician.de/software/pycuda/>
- <http://mathematician.de/software/pyopencl/>