

# Lecture 1: Introduction

Parallel Programming (INF-3201)  
University of Tromsø

John Markus Bjørndalen



1

## Instructors

		Office	Email
John Markus Bjørndalen	Instructor	REALF A247	<a href="mailto:jmb@cs.uit.no">jmb@cs.uit.no</a>
Edvard Pedersen	Assistant		<a href="mailto:gmtired@gmail.com">gmtired@gmail.com</a>

2

## What is Parallel Programming?

### Sequential programming:

- Instructions executed one by one
- Only one at any point in time




Figure from Blaise Barney, "Introduction to Parallel Computing", Livermore Computing.

3

## What is Parallel Programming?

### Parallel programming

- Break program into smaller parts
- Execute instructions from each part in parallel




Figure from Blaise Barney, "Introduction to Parallel Computing", Livermore Computing.


4

## Why use Parallel Programming?

### Save time and money

Parallel clusters from cheap, commodity components.

GPUs (many cores) can solve problems in real-time that we can't compute fast enough on a single core.



<http://top500.org/statistics/vertime/>

5

## Why use Parallel Programming?

### Solve larger problems

- Impractical to solve on a single computer
- Limited
  - Memory
  - CPU
  - Data/storage



[https://en.wikipedia.org/wiki/Supercomputer#/media/File:IBM\\_Blue\\_Gene\\_P\\_supercomputer.jpg](https://en.wikipedia.org/wiki/Supercomputer#/media/File:IBM_Blue_Gene_P_supercomputer.jpg)


6

## Why use Parallel Programming?

Solve larger problems

Examples:

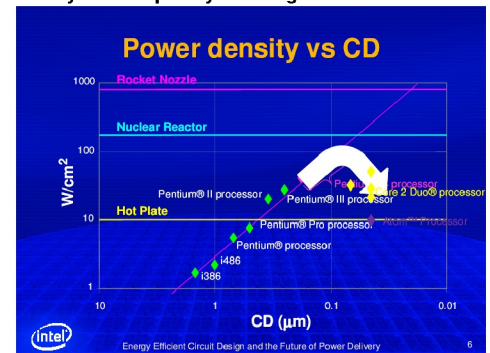
- Search engines
- Big data
- Weather forecasting
- Particle research



7

## Why use Parallel Programming?

Chip density and frequency: no longer «free lunch»



Energy Efficient Circuit Design and the Future of Power Delivery, Greg Taylor, EPEPS 2009

8


## Scaling after hitting the wall

Clock speed hits a wall due to fundamental physics

Frequencies beyond 5GHz melt chips

The free lunch is over

Moore's Law: transistor density doubles every 18 months



9

## Scaling after hitting the wall

How to increase chip performance?

Chip density is still increasing

⇒ Can increase #cores/chip to increase chip performance and reduce power




Image from <http://www.proportional.com/2012/08/01/e-close-look-at-xeon-phi-intel-50-core-best/>

10

## Goals of the course

- Provide students with theory and practice on parallel programming
  - **Theory**
    - Fundamental concepts of parallel programming
    - Issues of parallel performance
    - Parallel programming techniques
    - Parallel algorithms
  - **Practice**
    - Parallel programming languages, environments and toolboxes
    - Hands-on experience with multi-core processors, graphics processors (GPU) and clusters.
    - Benchmarking and analyzing parallel programs

11

## Theory


- Fundamental concepts of parallel programming
  - Parallel computers
    - message-passing multicomputer, shared memory multiprocessor, classification (Flynn's taxonomy), ...
  - Parallel programming models
    - message-passing (**clusters**), shared memory (**Intel multicore CPU**), data-parallel model (**Nvidia GPU**), ...
- Issues of parallel performance
  - Parallelism (Amdahl's law, DAG model), race-conditions, deadlocks, ...
- Parallel programming techniques
  - Divide-and-conquer, load balancing, data-parallel exploitation, ...
- Parallel algorithms
  - Sorting, numerical algorithms, prefix sums, ...

12

## Example: Matrix multiplication

### Sequential program

```
for i = 1 to n do
  for j = 1 to n do
    x[i,j] = 0
    for k = 1 to n do
      x[i,j] = x[i,j] + A[i,k] * B[k,j]
    ⇒ execution time  $O(n^3)$ 
```



$$X_{n,n} = A_{n,n} \cdot B_{n,n}$$

$$x_{1,2} = a_{1,1}b_{1,2} + a_{1,2}b_{2,2}$$


$$x_{3,3} = a_{3,1}b_{1,3} + a_{3,2}b_{2,3}$$

13

## Example: Matrix multiplication

### Sequential program

```
for i = 1 to n do
  for j = 1 to n do
    x[i,j] = 0
    for k = 1 to n do
      x[i,j] = x[i,j] + A[i,k] * B[k,j]
    ⇒ execution time  $O(n^3)$ 
```



### Parallel program

Different  $X_{i,j}$ s are independent

⇒ 2 outer loops can be parallelized

⇒ execution time  $O(n)$

(with some assumptions – can you spot them?)

$$X_{n,n} = A_{n,n} \cdot B_{n,n}$$


$$x_{1,2} = a_{1,1}b_{1,2} + a_{1,2}b_{2,2}$$

$$x_{3,3} = a_{3,1}b_{1,3} + a_{3,2}b_{2,3}$$

14

## Practice

- Parallel programming languages, environments and toolboxes
  - Message-passing model: [Open MPI](#), an interface (API) & tools for developing large scale message-passing applications
  - Shared memory model: [OpenMP](#), pragma based extension to C for improving performance on multi-core processors
  - Data-parallel model: [CUDA/OpenCL](#), environments & tools for leveraging the massively parallel processing power of GPUs



NB: we will be making some changes to the projects this year. You will get more information later.

15

## Projects

- You are given sequential codes
  - Parallelize code
  - Execute on cluster, multi-core computers and graphics processors (GPUs)
  - Evaluate results (speedup vs. expected)
  - Explain WHY!
- 3 mandatory (P1,2,3)

## Course schedule

Monday	Friday
14:15 - 16:00 Lecture, REALF B203	08:15 – 10:00 Colloquium, REALF A016
	10:15-11:00 Lecture, REALF B203

### Assessment

- 3 mandatory assignments (pass/fail)
- 1 final written exam: 100%

17

## Important messages

- Make sure you have
  - Access to the course room in Fronter <https://fronter.com/uit/>
    - Schedules, handouts, important messages, hand-ins,...
    - Note that plans are tentative
  - Access to the gitlab repository:
    - <https://source.uit.no/ifi-courses/inf-3201-2015-resources/tree/master>
    - it should be public/open, so you don't need to log in
    - There is a link to the gitlab repository in the Fronter room (under archives).

18

## References

---

- Blaise Barney, "Introduction to Parallel Computing", Livermore Computing.
- Charles Leiserson et al., "How to Survive the Multicore Revolution", Cilk Arts.
- Saman Amarasinghe, 6.189 "Multicore Programming Primer", January (IAP) 2007. (MIT OpenCourseWare). <http://ocw.mit.edu> (accessed 21 07, 2010). License: Creative Commons Attribution-Noncommercial-Share Alike.
- Guy Blelloch & Bruce Maggs. "Parallel Algorithms" (from Computer Science Handbook, 2<sup>nd</sup> Edition, ISBN-13: 978-1594933609)