

## Lecture 2: Parallel Computers

Parallell Programming (INF-3201)


John Markus Bjørndalen



1

## Speedup Factor

$$S(p) = \frac{\text{Execution time using one processor (best sequential algorithm)}}{\text{Execution time using multiprocessor with } p \text{ processors}} = \frac{t_s}{t_p}$$



$S(p)$  : increase in speed by using multiprocessor.

Use best sequential algorithm with single processor system.

Underlying algorithm for parallel implementation is usually different.

2

## Speedup Factor


Can also be computed in terms of computational steps as (sequential and parallel) time complexity.

Ignore communication overhead in parallel implementation

$$S(p) = \frac{\text{Number of computational steps using one processor}}{\text{Number of parallel computational steps with } p \text{ processors}}$$

3

## Maximum Speedup - Amdahl's law



Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 3rd Edition, by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

4

## Maximum Speedup - Amdahl's law


Speedup factor is given by:

$$S(p) = \frac{t_s}{f t_s + (1-f) t_s / p} = \frac{p}{1 + (p-1)f}$$

This equation is known as Amdahl's law

5

## Speedup against number of processors



Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 3rd Edition, by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

6

## Maximum speedup

$$S(p) = \frac{t_s}{ft_s + (1-f)t_s/p} = \frac{p}{1 + (p-1)f}$$

$$S(p)_{p \rightarrow \infty} = \frac{1}{f}$$


Example:  $f = 5\% \Rightarrow \text{maximum } S(p) = 20$

7

## Maximum Speedup

Maximum speedup: usually  $p$  with  $p$  processors (linear speedup).

Is it possible to get superlinear speedup (greater than  $p$ )?




- Yes, but usually a specific reason such as:
  - Extra memory in multiprocessor system
  - Nondeterministic algorithm

8

## Superlinear Speedup example - Searching


Sequential:



9

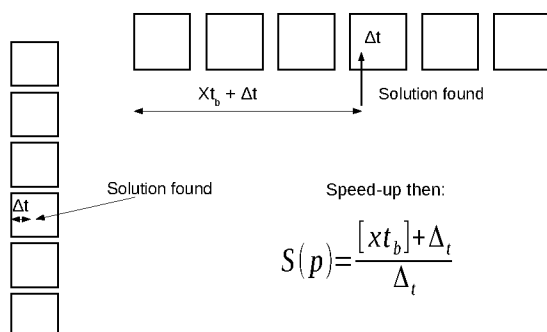
## Superlinear Speedup example - Searching

Parallel



10

## Superlinear Speedup example - Searching



NB: the book uses  $t_b = t_s/p$

11

## Efficiency

- How efficiently processors are being used on the computation (instead of idling, overheads etc)

$$E = \frac{\text{Execution time using one processor}}{\text{Execution time using a multiprocessor} \times \text{number of processors}}$$


$$E = \frac{t_s}{t_p \times p} = \frac{S(p)}{p} \times 100\%$$

- $E = 100\% \Leftrightarrow S(p) = p$

12

## Conventional Computer

Consists of a processor executing a program stored in a (main) memory:




Each main memory location located by its address. Addresses start at 0 and extend to  $2^b - 1$  when there are  $b$  bits (binary digits) in address.

13


## Shared Memory Multiprocessor System

Natural way to extend single processor model - have multiple processors connected to multiple memory modules, such that each processor can access any memory module :



14


## Simplistic view of a small shared memory multiprocessor



15

## Quad Pentium Shared Memory Multiprocessor

NB: not quad-core!



16

## Heterogenous architectures

Example: IBM Cell Architecture

Used in Playstations and some supercomputing areas.





Figure from <http://arstechnica.com/features/2005/02/cell-2/>

## Distributed Shared Memory

Making main memory of group of interconnected computers look like a single memory with single address space. Then can use shared memory programming techniques.



18

## Programming Shared Memory Multiprocessors


- **Threads** - programmer decomposes program into individual parallel sequences, (threads), each being able to access variables declared outside threads.
  - Example: Pthreads
- Sequential programming language with **added syntax** to declare shared variables and specify parallelism.
  - Example: Intel Cilk++
- Sequential programming language with **preprocessor compiler directives** to declare shared variables and specify parallelism.
  - Example: OpenMP

## Programming Shared Memory Multiprocessors

- **Languages** – languages especially designed for concurrency or parallel programming
  - Example: Occam-pi, Go
- **Libraries**
  - Example: Intel TBB (template library), new concurrency mechanisms in Java

## Message-Passing Multicomputer

Complete computers connected through an interconnection network:



## Key concepts and issues


- Concepts
  - **Bandwidth**: data rate (bits/sec)
  - **Throughput**: average rate successfully transmitted.
  - **Latency**:
    - **Network latency**: time to make a message transfer through the network
    - **Communication latency**: total time to send a message, including software overhead and interface delays
    - **Message latency**: time to send a **zero-length** message
  - **Diameter**: minimum number of links between the two farthest nodes in the network
    - ⇒ Maximum distance a message must travel ⇒ network latency
- Issues in network design:
  - Bandwidth, latency and cost.

## Interconnection Networks

- With **direct links** between computers/processors
  - Exhaustive interconnections
  - 2- and 3-dimensional meshes
  - Hypercube
- Using **switches**:
  - Crossbar
  - Trees
  - Multistage interconnection networks


## Exhaustive Interconnections

- connect every computer to every other computers with links
    - fully connected network
  - Latency: diameter =
  - Cost: # links for  $c$  computers ?
    - $c(c-1)/2$
- ⇒ suitable only for a very small system




## Two-dimensional array (mesh)

For  $p^2$  processors, diameter =  $2(p-1)$   
 Also three-dimensional - used in some large high performance systems.



25


## Three-dimensional hypercube



For  $p$  processors, diameter =  $\log_2 p$

26


## 3D-hypercube or 3D-mesh ?



For  $p^3$  processors, diameter =  $3(p-1)$


27

## Four-dimensional hypercube




28

## Crossbar switch



29

## Tree

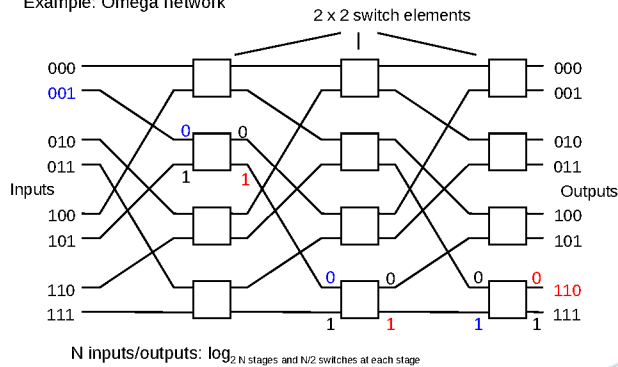


- Complete binary tree
- $p$  processors  $\Rightarrow$  height =  $\log_2 p$

30

## Multistage Interconnection Network

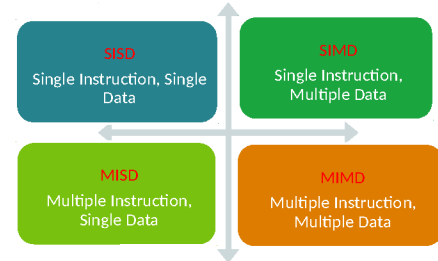
Example: Omega network



31

## Flynn's Classifications


- Flynn (1966) created a classification for computers based on the two dimensions of **Instruction** and **Data**.
  - Each dimension has two possible states: **Single** and **Multiple**.
- ⇒ 4 possible classifications



32

## Single Instruction, Single Data (SISD)


- A serial (non-parallel) computer
- Single instruction: only one instruction stream is executed by the CPU during one clock cycle
- Single data: only one data stream is used as input during one clock cycle
- Deterministic execution
- The oldest and the most common type of (single-core) computer



33

## Single Instruction, Multiple Data (SIMD):


- A type of parallel computer
- Single instruction: All processing units execute the same instruction at any given clock cycle
- Multiple data: Each processing unit can operate on a different data element
- Synchronous (lockstep) and deterministic execution
- Best suited for graphics/image processing.
  - Graphics Processing Units (GPU)



34

## Multiple Instruction, Single Data (MISD)


- A single data stream is fed into multiple processing units.
- Each processing unit operates on the data independently via independent instruction streams.
- Few actual examples of this class have ever existed.
- Some conceivable uses might be:
  - multiple frequency filters operating on a single signal stream
  - multiple cryptography algorithms attempting to crack a single coded message.



35

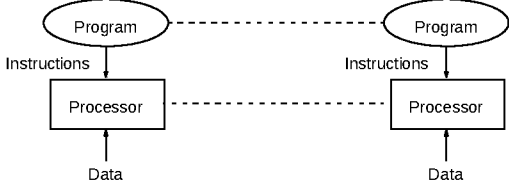
## Multiple Instruction, Multiple Data (MIMD)

- Currently, the most common type of parallel computer.
- Multiple Instruction: every processor may be executing a different instruction stream
- Multiple Data: every processor may be working with a different data stream
- Examples: most current supercomputers, clusters, multi-core computers.
- Note: many MIMD architectures also include SIMD sub-components



36

## Other classifications

- Multiple **Program** Multiple Data (MPMD)
    - Each processor will have its own program to execute
- 
- Single **Program** Multiple Data (SPMD)
    - Single source program written and each processor executes its personal copy of this program
    - Parts of the program can be executed by certain processors and not others depending on processor ID.

37

## Networked Computers as a Computing Platform

- A network of computers became a very attractive alternative to expensive supercomputers and parallel computer systems for high-performance computing in early 1990's.
- Several early projects. Notable:
  - Berkeley NOW (network of workstations) project.
  - NASA Beowulf project.

38

## Key advantages:

- Very high performance workstations and PCs readily available at low cost.
- The latest processors can easily be incorporated into the system as they become available.
- Existing software can be used or modified.

39

## Software Tools for Clusters

- Based upon Message Passing Parallel Programming:
- Parallel Virtual Machine (PVM) - developed in late 1980's. Became very popular.
- Message-Passing Interface (MPI) - standard defined in 1990s.
- Both provide a set of user-level libraries for message passing. Use with regular programming languages (C, C++, ...)

40

## Cluster Interconnects

- Originally fast Ethernet on low cost clusters
- Gigabit Ethernet - easy upgrade path

### More Specialized/Higher Performance

- Myrinet - 2.4 Gbits/sec
- InfiniBand
- SCI (Scalable Coherent Interface)
- QsNet
- ...

41

## References

- Blaise Barney, "Introduction to Parallel Computing", Livermore Computing.
- Barry Wilkinson & Michael Allen. Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers.

42