


Lecture 3b: Message-Passing Computing - MPI

Parallell Programming (INF-3201)

John Markus Bjørndalen

MPI (Message Passing Interface)

- What?
 - a **specification** for the developers and users of message passing libraries (not a library).
- Why use MPI?
 - **Standardization** - MPI is the only message passing library interface which can be considered a standard.
 - **Portability** - There is no need to modify your source code when you port your application to a different platform.
 - **Performance** - Vendor implementations can exploit native hardware features to optimize performance.
 - **Functionality** - Over 115 routines are defined in MPI-1 alone.
 - **Availability** - A variety of implementations are available.



MPI Process Creation and Execution

- Purposely not defined - Will depend upon implementation.
- MPI version 1
 - Only static process creation supported.
 - All processes must be defined prior to execution and started together.
- MPI version 2
 - Dynamic process creation supported (MPI_Comm_spawn)
- Originally SPMD model of computation.
 - MPMD also possible with static creation


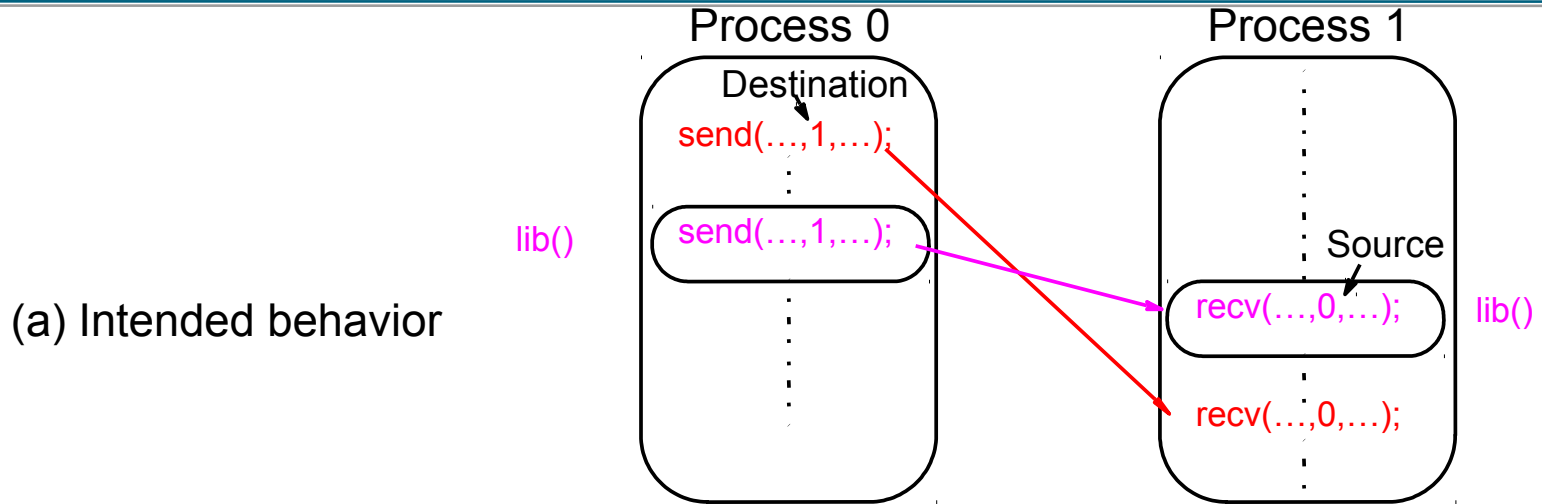
Using SPMD Computational Model

```
main (int argc, char *argv[])
{
    MPI_Init(&argc, &argv);
    .
    .
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank); /*find process rank */

    if (myrank == 0)
        master();
    else
        slave();
    .
    .
    MPI_Finalize();
}
```

where master() and slave() are to be executed by master process and slave process, respectively.


Unsafe message passing - Example



MPI Solution: “Communicators”

- Defines a communication domain - a set of processes that are allowed to communicate between themselves.
- Processes have ranks associated with communicator (0 to n-1)
- Initially, all processes enrolled in a “universe” called `MPI_COMM_WORLD`
- Other communicators can be established for groups of processes.
- Communication domains of libraries can be separated from that of a user program.
- Used in all point-to-point and collective MPI message-passing communications.

Communicator Example



MPI Point-to-Point Communication

- Involves one source process and one destination process
- Uses send and receive routines with message tags (and communicator).
- Wild card message tags available

MPI Blocking Routines

- Return when **locally complete**
 - location used to hold message can be used again or altered without affecting message being sent.
- When **blocking send** returns
 - process free to move on without adversely affecting message.
 - does not mean that message has been received,

Parameters of blocking send

MPI_Send(buf, count, datatype, dest, tag, comm)

Address of
send buffer

Number of items
to send

Datatype of
each item

Rank of destination
process

Message tag

Communicator

Parameters of blocking receive

MPI_Recv(buf, count, datatype, src, tag, comm, status)

Address of
receive buffer

Maximum number
of items to receive

Datatype of
each item

Rank of source
process

Message tag

Communicator

Status
after operation

Example

To send an integer x from process 0 to process 1,

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank); /* find rank */

if (myrank == 0) {
    int x;
    MPI_Send(&x, 1, MPI_INT, 1, msgtag, MPI_COMM_WORLD);
} else if (myrank == 1) {
    int x;
    MPI_Recv(&x, 1, MPI_INT, 0, msgtag, MPI_COMM_WORLD, status);
}
```

MPI Nonblocking Routines

- Nonblocking send - `MPI_Isend()` - will return immediately even before source location is safe to be altered.
- Nonblocking receive - `MPI_Irecv()` - will return even if no message to accept.

Nonblocking Routine Formats

```
MPI_Isend(buf, count, datatype, dest, tag, comm, request);
```

```
MPI_Irecv(buf, count, datatype, source, tag, comm, request);
```

Completion detected by `MPI_Wait()` and `MPI_Test()`.

```
MPI_Wait(request, ...)
```

waits until operation completed and returns then.

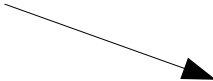
```
MPI_Test(request, ...)
```

returns with flag (set) indicating whether operation completed at that time.

Example

To send an integer x from process 0 to process 1 and allow process 0 to continue,

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank); /* find rank */
if (myrank == 0) {
    int x;
    MPI_Isend(&x,1,MPI_INT, 1, msgtag, MPI_COMM_WORLD, req1);
    compute();
    MPI_Wait(req1, status);
} else if (myrank == 1) {
    int x;
    MPI_Recv(&x,1,MPI_INT,0,msgtag, MPI_COMM_WORLD, status);
}
```



Send Communication Modes

- **Buffered Mode** - Send may start and return before a matching receive. Necessary to specify buffer space via routine `MPI_Buffer_attach()`.
- **Synchronous Mode** - Send and receive can start before each other but can only complete together.
- **Standard Mode** - Not assumed that corresponding receive routine has started. If buffering provided, send could complete before receive reached.
- **Ready Mode** - Send can only start if matching receive already reached, otherwise error. Use with care.

-
- Each of the four modes can be applied to both blocking and nonblocking send routines.
 - Any type of send routine can be used with any type of receive routine.

Collective Communication


- All or None
 - Must involves **all** processes in the scope of an intra-communicator.
 - It is the programmer's responsibility to insure that **all** processes within a communicator participate in any collective routines.
- Programming Considerations and Restrictions
 - Routines are blocking.
 - Routines do not take message tag arguments.
 - Can only be used with MPI predefined datatypes. Solution?
 - Convert user-defined datatypes to an array of integer (cf. cast operator in C)
- Principal collective routines:
 - **MPI_Bcast()** - Broadcast from root to all other processes
 - **MPI_Gather()** - Gather values for group of processes
 - **MPI_Scatter()** - Scatters buffer in parts to group of processes
 - **MPI_Alltoall()** - Sends data from all processes to all processes
 - **MPI_Reduce()** - Combine values on all processes to single value
 - **MPI_Reduce_scatter()** - Combine values and scatter results
 - **MPI_Scan()** - Compute prefix reductions of data on processes

MPI_Bcast

MPI_Bcast

Broadcasts a message to all other processes of that group

```
count = 1;  
source = 1;          broadcast originates in task 1  
MPI_Bcast(&msg, count, MPI_INT, source, MPI_COMM_WORLD);
```




MPI_Gather

MPI_Gather

Gathers together values from a group of processes

```
sendcnt = 1;  
recvcnt = 1;  
src = 1;           messages will be gathered in task 1  
MPI_Gather(sendbuf, sendcnt, MPI_INT,  
           recvbuf, recvcnt, MPI_INT,  
           src, MPI_COMM_WORLD);
```




MPI_Scatter

MPI_Scatter

Sends data from one task to all other tasks in a group

```
sendcnt = 1;  
recvcnt = 1;  
src = 1;           task 1 contains the message to be scattered  
MPI_Scatter(sendbuf, sendcnt, MPI_INT,  
            recvbuf, recvcnt, MPI_INT,  
            src, MPI_COMM_WORLD);
```



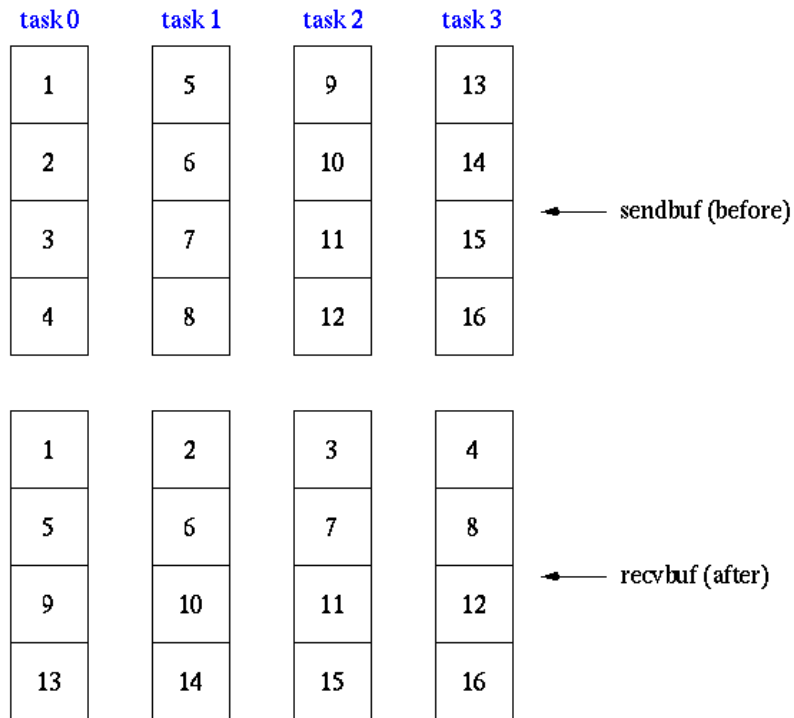
MPI_Alltoall

MPI_Alltoall

Sends data from all to all processes. Each process performs a scatter operation.

```
sendcnt = 1;  
recvcnt = 1;
```

```
MPI_Alltoall(sendbuf, sendcnt, MPI_INT,  
             recvbuf, recvcnt, MPI_INT,  
             MPI_COMM_WORLD);
```



MPI_Reduce

MPI_Reduce

Perform and associate reduction operation across all tasks in the group and place the result in one task

```
count = 1;
```

```
dest = 1;
```

result will be placed in task 1

```
MPI_Reduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,  
           dest, MPI_COMM_WORLD);
```

task 0

task 1

task 2

task 3

1

2

3

4

← sendbuf (before)

10

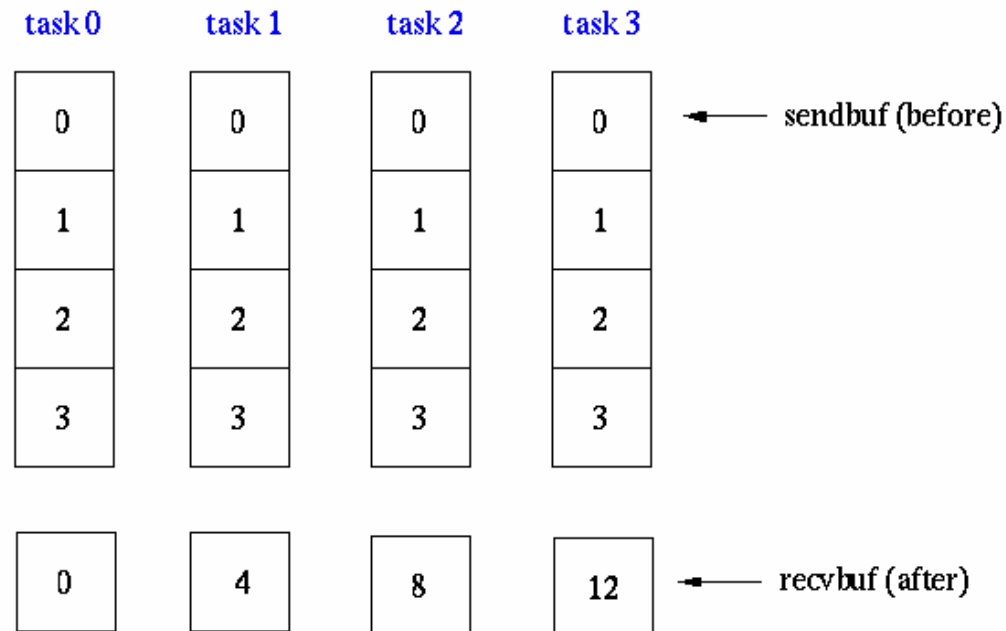
← recvbuf (after)

MPI_Reduce_scatter

MPI_Reduce_scatter

Perform reduction operation on vector elements across all tasks in the group, then distribute segments of result vector to tasks

```
recvcount = 1;  
MPI_Reduce_scatter(sendbuf, recvbuf, recvcount, MPI_INT, MPI_SUM,  
MPI_COMM_WORLD);
```




MPI_Scan

MPI_Scan

Computes the scan (partial reductions) of data on a collection of processes

```
count = 1;  
MPI_Scan(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,  
        MPI_COMM_WORLD);
```



Example

To gather items from group of processes into process 0, using dynamically allocated memory in root process:

```
int data[10];          /*data to be gathered from processes*/
MPI_Comm_rank(MPI_COMM_WORLD, &myrank); /* find rank */
if (myrank == 0) {
    MPI_Comm_size(MPI_COMM_WORLD, &grp_size); /*find group size*/
    buf = (int*)malloc(grp_size * 10 * sizeof(int)); /*allocate memory*/
}
MPI_Gather(data, 10, MPI_INT, buf, grp_size*10, MPI_INT, 0, MPI_COMM_WORLD);
```



gathers from all processes, including root.

Barrier routine

- A means of synchronizing processes by stopping each one until they all have reached a specific “barrier” call.
 - **MPI_Barrier (comm)**

Sample MPI program

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
#define MAXSIZE 1000
void main(int argc, char *argv)
{
    int myid, numprocs;
    int data[MAXSIZE], i, x, low, high, myresult, result;
    char fn[255];
    char *fp;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    if (myid == 0) { /* Open input file and initialize data */
        strcpy(fn,getenv("HOME"));
        strcat(fn,"/MPI/rand_data.txt");
        if ((fp = fopen(fn,"r")) == NULL) {
            printf("Can't open the input file: %s\n\n", fn);
            exit(1);
        }
        for(i = 0; i < MAXSIZE; i++) fscanf(fp,"%d", &data[i]);
    }
    MPI_Bcast(data, MAXSIZE, MPI_INT, 0, MPI_COMM_WORLD); /* broadcast data */
    x = n/nproc; /* Add my portion Of data */
    low = myid * x;
    high = low + x;
    for(i = low; i < high; i++)
        myresult += data[i];
    printf("I got %d from %d\n", myresult, myid); /* Compute global sum */
    MPI_Reduce(&myresult, &result, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
    if (myid == 0) printf("The sum is %d.\n", result);
    MPI_Finalize();
}
```

Estimating speedup

Sequential execution time, t_s : Estimate by counting computational steps of best sequential algorithm.

Parallel execution time, t_p : In addition to number of computational steps, t_{comp} , need to estimate communication overhead, t_{comm} :

$$t_p = t_{\text{comp}} + t_{\text{comm}}$$

Computation Time

- Count number of computational steps.
- When more than one process executed simultaneously, count computational steps of most complex process.
- Generally, function of number of data elements n and number of processors p , i.e.

$$t_{\text{comp}} = f(n, p)$$

Often break down computation time into parts. Then

$$t_{\text{comp}} = t_{\text{comp1}} + t_{\text{comp2}} + t_{\text{comp3}} + \dots$$

Analysis usually done assuming that all processors are same and operating at same speed.

Communication Time


Many factors, including network structure and network contention. As a first approximation, use

$$t_{\text{comm1}} = t_{\text{startup}} + nt_{\text{data}}$$

t_{startup} is startup time, essentially time to send a message with no data. Assumed to be constant.

t_{data} is transmission time to send one data word, also assumed constant, and there are n data words.

Idealized Communication Time



Final communication time, t_{comm}

Sum of communication times of all **sequential messages** from **a process**, i.e.

$$t_{\text{comm}} = t_{\text{comm1}} + t_{\text{comm2}} + t_{\text{comm3}} + \dots$$

- Assumption: Communication patterns of all processes are the same and take place together
⇒ only one process need be considered.

Both t_{startup} and t_{data} , measured in units of one computational step, so that can add t_{comp} and t_{comm} together to obtain parallel execution time, t_p .

Benchmark Factors

With t_s , t_{comp} , and t_{comm} , can establish speedup factor and computation/communication ratio for a particular algorithm/implementation.

$$\text{Speedup factor} = \frac{t_s}{t_p} = \frac{t_s}{t_{comp} + t_{comm}}$$

$$\text{Computation/communication ratio} = \frac{t_{comp}}{t_{comm}}$$

Both functions of number of processors, p , and number of data elements, n .

Factors give indication of **scalability** of parallel solution with increasing number of processors and problem size.

Computation/communication ratio will highlight **effect of communication** with increasing problem size and system size.

References

- Barry Wilkinson & Michael Allen. Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers.
- Blaise Barney, “Message Passing Interface (MPI)”, Livermore Computing.