

## Lecture 8: Synchronous Computations

Parallel Programming (INF-3201)  
Autumn 2013

John Markus Bjørndalen




1

## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

2


## Processes reaching barrier at different times



3

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

In message-passing systems, barriers provided with library routines:



**MPI**  
**MPI\_Barrier()**

Barrier with a named communicator being the only parameter.


Called by each process in the group, blocking until all members of the group have reached the barrier call and only returning then.

4

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Barrier Implementation

Centralized counter implementation (a linear barrier):




5

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Barrier Implementation

Centralized counter implementation (a linear barrier):



Good barrier implementations must take into account that a barrier might be used more than once in a process.


Might be possible for a process to enter the barrier for a second time before previous processes have left the barrier for the first time.

6

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Barrier Implementation


Centralized counter implementation (a linear barrier):



7

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Barrier implementation in a message-passing system



- Master:**

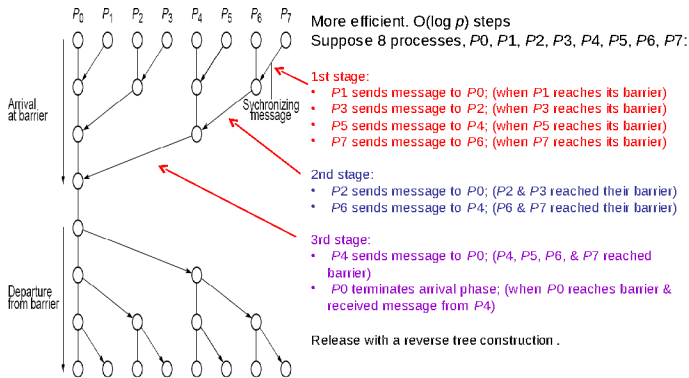
```
for (i = 0; i < n; i++) /* count slaves as they reach barrier */
    recv(P_slave);
for (i = 0; i < n; i++) /* release slaves */
    send(P_i);
```
- Slave processes:**

```
send(P_master);
recv(P_master);
```

8

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Tree Implementation



9

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.


## Butterfly Barrier

- 1st stage  $P_0 \leftrightarrow P_1, P_2 \leftrightarrow P_3, P_4 \leftrightarrow P_5, P_6 \leftrightarrow P_7$   
 2nd stage  $P_0 \leftrightarrow P_2, P_1 \leftrightarrow P_3, P_4 \leftrightarrow P_6, P_5 \leftrightarrow P_7$   
 3rd stage  $P_0 \leftrightarrow P_4, P_1 \leftrightarrow P_5, P_2 \leftrightarrow P_6, P_3 \leftrightarrow P_7$

Textbook: At stage  $s$ ,  $p_i$  synchronizes with  $p_j$ , where  $j = i + 2s - 1$ . Wrong!

Correct:  $j = i \text{ XOR } 2s - 1$

(cf. D. E. Culler et al. Parallel Computer Architecture – A Hardware/Software Approach). jPhuong




10

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Local Synchronization

Suppose a process  $P_i$  needs to be synchronized and to exchange data with process  $P_{i-1}$  and process  $P_{i+1}$  before continuing:



Not a perfect three-process barrier because process  $P_{i-1}$  will only synchronize with  $P_i$  and continue as soon as  $P_i$  allows. Similarly, process  $P_{i+1}$  only synchronizes with  $P_i$ .

11

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Deadlock

When a pair of processes each send and receive from each other, deadlock may occur.

Deadlock will occur if both processes perform the **send**, using **synchronous routines** first (or blocking routines without sufficient buffering). This is because neither will return; they will wait for matching receives that are never reached.

12

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Deadlock conditions

Necessary conditions for deadlocks:

- **Mutual exclusion**: At least two resources must be non-shareable
- **Hold and Wait / Resource holding**: processes can hold a resource while requesting another
- **No Preemption**: resources must be released voluntarily – cannot force a process to release it (or take it away)
- **Circular wait**: a set of processes form a circle, where each wait for a resource held by the next in the circle.

13

## A Solution

Arrange for one process to receive first and then send and the other process to send first and then receive.

### Example

Linear pipeline, deadlock can be avoided by arranging so the even-numbered processes perform their sends first and the odd-numbered processes perform their receives first.

14

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Combined deadlock-free blocking sendrecv() routines

### Example

```
Process Pi-1      Process Pi      Process Pi+1  
sendrecv(Ri) ; ↔ sendrecv(Ri-1) ;  
                  sendrecv(Ri+1) ; ↔ sendrecv(Ri) ;
```

MPI provides `MPI_Sendrecv()` and `MPI_Sendrecv_replace()`.  
**MPI sendrecv()s actually has 12 parameters!**

15

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

16

## Synchronized Computations

Can be classified as:

- Fully synchronous
- or
- Locally synchronous

In fully synchronous, all processes involved in the computation must be synchronized.

In locally synchronous, processes only need to synchronize with a set of logically nearby processes, not all processes involved in the computation

17

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Fully Synchronized Computation Examples

### Data Parallel Computations

Same operation performed on different data elements simultaneously; i.e., in parallel.

Particularly convenient because:

- Ease of programming (essentially only one program).
- Can scale easily to large problem sizes. How large?
- Many numeric and some non-numeric problems can be cast in a data parallel form.

18

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Example

To add the same constant to each element of an array:

```
for (i = 0; i < n; i++)  
    a[i] = a[i] + k;
```

The statement:


```
a[i] = a[i] + k;
```

could be executed simultaneously by multiple processors, each using a different index  $i$  ( $0 < i \leq n$ ).

19

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Data Parallel Computation



20

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## forall construct

Special "parallel" construct in parallel programming languages to specify data parallel operations

### Example

```
forall (i = 0; i < n; i++) {  
    body  
}
```

states that  $n$  instances of the statements of the body can be executed simultaneously.

One value of the loop variable  $i$  is valid in each instance of the body, the first instance has  $i = 0$ , the next  $i = 1$ , and so on.

21

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

To add  $k$  to each element of an array,  $a$ , we can write

```
forall (i = 0; i < n; i++)  
    a[i] = a[i] + k;
```

22

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

Data parallel technique applied to multiprocessors and multicomputers

### Example

To add  $k$  to the elements of an array:

```
i = myrank;  
a[i] = a[i] + k;                      /* body */  
barrier(mygroup);
```

where  $\text{myrank}$  is a process rank between 0 and  $n - 1$ .

23

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

24

## Data Parallel Example Prefix Sum Problem

Given a list of numbers,  $x_0, \dots, x_{n-1}$ , compute all the partial summations, i.e.:

$$x_0 + x_1; x_0 + x_1 + x_2; x_0 + x_1 + x_2 + x_3; \dots$$


Can also be defined with associative operations other than addition.

Widely studied. Practical applications in areas such as processor allocation, data compaction, sorting, and polynomial evaluation.

25

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Data parallel prefix sum operation



26

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

### Sequential code

```
for (j = 0; j < log(n); j++) /* at each step, add */
  for (i = 2j; i < n; i++) /* to accumulating sum */
    x[i] = x[i] + x[i - 2j];
```

### Parallel code

```
for (j = 0; j < log(n); j++) /* at each step, add */
  forall (i = 0; i < n; i++) /* to sum */
    if (i >= 2j) x[i] = x[i] + x[i - 2j];
```

27

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

28

## Synchronous Iteration (Synchronous Parallelism)

Each iteration composed of several processes that start together at beginning of iteration. Next iteration cannot begin until all processes have finished previous iteration.

Using **forall** construct:

```
for (j = 0; j < n; j++) {           // for each synch. iteration
  forall (i = 0; i < N; i++) {       // N procs each using
    body(i);                         // specific value of i
  }
}
```

29

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

Using message passing barrier:

```
for (j = 0; j < n; j++) {           /*for each synchr.iteration */
  i = myrank;                       /*find value of i to be used */
  body(i);
  barrier(mygroup);
}
```

30

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Another fully synchronous computation example

**Solving a General System of Linear Equations by Iteration**  
Suppose the equations are of a general form with  $n$  equations and  $n$  unknowns

$$a_{n-1,0}x_0 + a_{n-1,1}x_1 + a_{n-1,2}x_2 \dots + a_{n-1,n-1}x_{n-1} = b_{n-1}$$

.

$$\begin{array}{rcl} a_{2,0}x_0 + a_{2,1}x_1 + a_{2,2}x_2 \dots & + a_{2,n-1}x_{n-1} & = b_2 \\ a_{1,0}x_0 + a_{1,1}x_1 + a_{1,2}x_2 \dots & + a_{1,n-1}x_{n-1} & = b_1 \\ a_{0,0}x_0 + a_{0,1}x_1 + a_{0,2}x_2 \dots & + a_{0,n-1}x_{n-1} & = b_0 \end{array}$$

where the unknowns are  $x_0, x_1, x_2, \dots, x_{n-1}$  ( $0 \leq i < n$ ).

31

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

By rearranging the  $i$ th equation:

$$a_{i,0}x_0 + a_{i,1}x_1 + a_{i,2}x_2 \dots + a_{i,n-1}x_{n-1} = b_i$$

to

$$x_i = (1/a_{i,i})[b_i - (a_{i,0}x_0 + a_{i,1}x_1 + a_{i,2}x_2 \dots + a_{i,i-1}x_{i-1} + a_{i,i+1}x_{i+1} \dots + a_{i,n-1}x_{n-1})]$$

or

$$x_i = \frac{1}{a_{i,i}} \left[ b_i - \sum_{j \neq i} a_{i,j} x_j \right]$$

This equation gives  $x_i$  in terms of the other unknowns.

Can be used as an iteration formula for each of the unknowns to obtain better **approximations**.

32

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Jacobi Iteration

All values of  $x$  are updated **together**.

Can be proven that Jacobi method will converge if diagonal values of  **$a$**  have an absolute value greater than sum of the absolute values of the other  **$a$** 's on the row (the array of  **$a$** 's is *diagonally dominant*) i.e. if

$$\sum_{j \neq i} |a_{i,j}| < |a_{i,i}|$$

This condition is a sufficient but not a necessary condition.

33

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Termination

A simple, common approach. Compare values computed in one iteration to values obtained from the previous iteration. Terminate computation when all values are within given tolerance; i.e., when

$$|x_i^t - x_i^{t-1}| < \text{error tolerance}$$


for all  $i$ , where  $x_i^t$  is the value of  $x_i$  after the  $t$ th iteration and  $x_i^{t-1}$  is the value of  $x_i$  after the  $(t-1)$ th iteration.

**However, this does not guarantee the solution to that accuracy.**

34

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Convergence Rate



35

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Parallel Code

Process  $P_i$  could be of the form

```
x[i] = b[i]; //initialize unknown
for (iteration = 0; iteration < limit; iteration++) {
    sum = -a[i][i] * x[i];
    for (j = 0; j < n; j++) // compute summation
        sum = sum + a[i][j] * x[j];
    new_x[i] = (b[i] - sum) / a[i][i]; // compute unknown
    allgather(&new_x[i], &x[i]); // bcst/rec values
    global_barrier(); // wait for all procs
}
```

**allgather()** sends the newly computed value of  **$x[i]$**  from process  $i$  to every other process and collects data broadcast from the other processes.


36

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

Introduce a new message-passing operation - Allgather.

## Allgather

Broadcast and gather values in one composite construction.



37

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Partitioning

Usually number of processors much fewer than number of data items to be processed. Partition the problem so that processors take on more than one data item.

38

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

**block allocation** – allocate groups of consecutive unknowns to processors in increasing order.

- processor  $P_0$  is allocated  $x_0, x_1, x_2, \dots, x_{(n/p)-1}$ ,
- processor  $P_1$  is allocated  $x_{n/p}, x_{n/p+1}, \dots, x_{(2n/p)-1}$ , and so on.

**cyclic allocation** – processors are allocated one unknown in order:

- processor  $P_0$  is allocated  $x_0, x_p, x_{2p}, \dots, x_{((n/p)-1)p}$ ,
- processor  $P_1$  is allocated  $x_1, x_{p+1}, x_{2p+1}, \dots, x_{((n/p)-1)p+1}$ , and so on.


39

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Effects of computation and communication in Jacobi iteration

Consequences of different numbers of processors done in textbook.

Get:



40

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

41

## Locally Synchronous Computation Heat Distribution Problem

An area has known temperatures along each of its edges.

Find the temperature distribution within.

42

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

# Heat Distribution Problem


Divide area into fine mesh of points,  $h_{ij}$ ,  $0 \leq i, j \leq n$

Temperature at an inside point taken to be average of temperatures of four neighboring points. Convenient to describe edges by points.

Temperature of each point by iterating the equation:

$$h_{i,j} = \frac{h_{i-1,j} + h_{i+1,j} + h_{i,j-1} + h_{i,j+1}}{4}$$

( $0 < i < n$ ,  $0 < j < n$ ) for a fixed number of iterations or until the difference between iterations less than some very small amount.



43

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

# Sequential Code

Using a fixed number of iterations

```
for (iteration = 0; iteration < limit; iteration++) {
    for (i = 1; i < n; i++)
        for (j = 1; j < n; j++)
            g[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
    for (i = 1; i < n; i++)
        for (j = 1; j < n; j++)
            h[i][j] = g[i][j];
}
```

using original numbering system ( $n \times n$  array).

44

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

To stop at some precision:

```
do {
    for (i = 1; i < n; i++)
        for (j = 1; j < n; j++)
            g[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);

    continue = FALSE;
    for (i = 1; i < n; i++)
        for (j = 1; j < n; j++)
            if (!converged(i,j)) {
                continue = TRUE;
                break;
            }

    for (i = 1; i < n; i++)
        for (j = 1; j < n; j++)
            h[i][j] = g[i][j];
} while (continue == TRUE);
```

45

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

# Parallel Code

With fixed number of iterations,  $P_{ij}$

(except for the boundary points):

```
for (iteration = 0; iteration < limit; iteration++) {
    g = 0.25 * (w + x + y + z);
    send(&g, P_{i-1,j}); /* non-blocking sends */
    send(&g, P_{i+1,j});
    send(&g, P_{i,j-1});
    send(&g, P_{i,j+1});
    recv(&w, P_{i-1,j}); /* synchronous receives */
    recv(&x, P_{i+1,j});
    recv(&y, P_{i,j-1});
    recv(&z, P_{i,j+1});
}
```

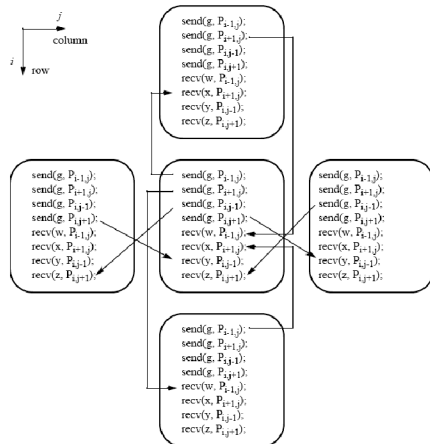
Local barrier

Important to use **send()**s that do not block while waiting for **recv()**s; otherwise processes would deadlock, each waiting for a **recv()** before moving on - **recv()**s must be synchronous and wait for **send()**s.

46

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

# Message passing for heat distribution problem



47

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

Version where processes stop when they reach their required precision:

```
iteration = 0;
do {
    iteration++;
    g = 0.25 * (w + x + y + z);
    send(&g, P_{i-1,j}); /* locally blocking sends */
    send(&g, P_{i+1,j});
    send(&g, P_{i,j-1});
    send(&g, P_{i,j+1});
    recv(&w, P_{i-1,j}); /* locally blocking receives */
    recv(&x, P_{i+1,j});
    recv(&y, P_{i,j-1});
    recv(&z, P_{i,j+1});
} while ((!converged(i, j)) || (iteration < limit));
send(&g, &i, &j, &iteration, &master);
```

48

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.



To handle the processes operating at the edges:

**MPI has a construct to help here**

```
if (last_row) x = bottom_value;
if (first_row) w = top_value;
if (first_column) y = left_value;
if (last_column) z = right_value;
iteration = 0;
do {
    iteration++;
    g = 0.25 * (w + x + y + z);
    if !(first_row) send(&g, R_{-1,j});
    if !(last_row) send(&g, R_{+1,j});
    if !(first_column) send(&g, R_{j,-1});
    if !(last_column) send(&g, R_{j,+1});
    if !(first_row) recv(&w, R_{-1,j});
    if !(last_row) recv(&x, R_{+1,j});
    if !(first_column) recv(&y, R_{j,-1});
    if !(last_column) recv(&z, R_{j,+1});
} while(!converged || (iteration < limit));
send(&g, &i, &j, iteration, R_{aster});
```

MPI provides Virtual Topologies to map processes into a mesh. /Phuong

49

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.


## Outline

- Synchronization
  - Barrier implementation
- Synchronized computations
  - Data parallel computations
    - Prefix Sum Problem
  - Synchronous iterations
    - General system of linear equations
  - Locally synchronous computations
    - Heat distribution
    - Implementation issues
      - Partitioning, deadlock

50

## Partitioning

Normally allocate more than one point to each processor, because many more points than processors.  
Points could be partitioned into square blocks or strips:




51

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Block partition

Four edges where data points exchanged.  
Communication time given by

$$t_{\text{commsq}} = 8 \left( t_{\text{startup}} + \frac{n}{\sqrt{p}} t_{\text{data}} \right)$$




52

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Strip partition

Two edges where data points are exchanged.  
Communication time is given by

$$t_{\text{commcol}} = 4(t_{\text{startup}} + nt_{\text{data}})$$



53

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Optimum

In general, strip partition best for large startup time, and block partition best for small startup time.

With the previous equations, block partition has a larger communication time than strip partition if

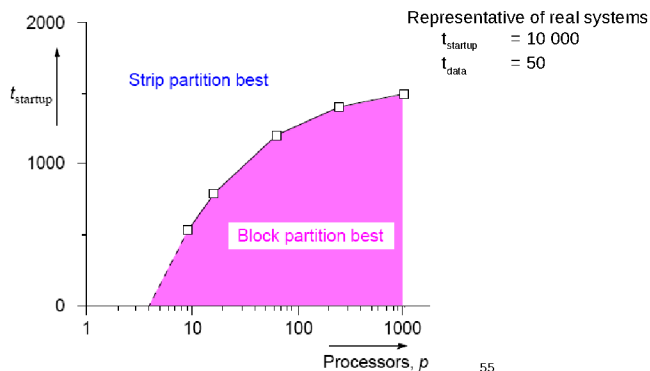
$$t_{\text{startup}} > n \left( 1 - \frac{2}{\sqrt{p}} \right) t_{\text{data}}$$

( $p \geq 9$ ).

54

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.


## Startup times for block and strip partitions



Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Ghost Points

Additional row of points at each edge that hold values from adjacent edge. Each array of points increased to accommodate ghost rows.



Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Safety and Deadlock

When all processes send their messages **first** and then receive all of their messages is **"unsafe"** because it relies upon buffering in the **send()**s. The amount of buffering is not specified in MPI.

If insufficient storage available, send routine may be delayed from returning until storage becomes available or until the message can be sent without buffering.

Then, a locally blocking **send()** could behave as a **synchronous send()**, only returning when the matching **recv()** is executed. Since a matching **recv()** would never be executed if all the **send()**s are synchronous, **deadlock would occur**.

57

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## Making the code safe

**Alternate the order** of the **send()**s and **recv()**s in adjacent processes so that only one process performs the **send()**s first.

Then even synchronous **send()**s would not cause deadlock.

**Good way you can test for safety is to replace message-passing routines in a program with synchronous versions.**

58

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## MPI Safe Message Passing Routines

MPI offers several methods for safe communication:

- **Combined send and receive routines:**  
`MPI_Sendrecv()`  
**which is guaranteed not to deadlock**
- **Buffered send(s):**  
`MPI_Bsend()`  
**here the user provides explicit storage space**
- **Nonblocking routines:**  
`MPI_Isend()` and `MPI_Irecv()`  
**which return immediately.**  
**Separate routine used to establish whether message has been received:**  
`MPI_Wait()`, `MPI_Waitall()`, `MPI_Waitany()`, `MPI_Test()`,  
`MPI_Testall()`, **or** `MPI_Testany()`.

59

Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen, © 2004 Pearson Education Inc. All rights reserved.

## References

- Barry Wilkinson & Michael Allen. Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers.
- D. E. Culler et al.. Parallel Computer Architecture – A Hardware/Software Approach